



Ai-Spy

IMAGE RECOGNITION

Project 2 Team:

Grace Ho

Pamela Moreno

Jay Richardson

Muhammad Hasnain Dosani



Project Goal

Train a machine learning model to
successfully recognise American Sign Language
digit and alphabet hand signals.



Methodology

Build initial Machine Learning model to successfully recognises ASL¹ digit hand signals. (10 symbols)

Determine if the model can be re-fit for both ASL¹ alphabet & digit hand signals. (34 symbols²)

Test additional hand pictures to see how model responds to other hand shapes.

¹ American Sign Language.

² Excluding characters J & Z which is movement based.

Digits Model

Dataset:

- 2062 images of ASL digit hand symbols
- 64 x 64 pixel black and white images, in numpy array (X & y files)
- 80/20 train vs test split



Digits Model cont.

Structure & compile used

- Keras Sequential model
- 16 layers: Convolutional 2D, MaxPooling2D, Drop out and Dense layers
- Relu activation, with Softmax final layer
- Adam optimiser, categorical_crossentropy loss

```
model.summary()
```


Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 64, 64, 8)	208
max_pooling2d (MaxPooling2D)	(None, 32, 32, 8)	0
dropout (Dropout)	(None, 32, 32, 8)	0
conv2d_1 (Conv2D)	(None, 32, 32, 16)	1168
max_pooling2d_1 (MaxPooling2D)	(None, 16, 16, 16)	0
dropout_1 (Dropout)	(None, 16, 16, 16)	0
conv2d_2 (Conv2D)	(None, 16, 16, 32)	4640
max_pooling2d_2 (MaxPooling2D)	(None, 8, 8, 32)	0
dropout_2 (Dropout)	(None, 8, 8, 32)	0
conv2d_3 (Conv2D)	(None, 8, 8, 64)	18496
max_pooling2d_3 (MaxPooling2D)	(None, 4, 4, 64)	0
dropout_3 (Dropout)	(None, 4, 4, 64)	0
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 128)	131200
dense_1 (Dense)	(None, 64)	8256
dense_2 (Dense)	(None, 10)	650
=====		
Total params: 164,618		
Trainable params: 164,618		
Non-trainable params: 0		

Digits Model cont. - CNN

A convolutional layer contains a set of parameters that need to be learned. Each filter is convolved with the input volume to compute an activation map.


The activation map is slid across the width and height of the input in a "pool size" set in the model architecture with the weights computed at every spatial position.




Digits Model cont.

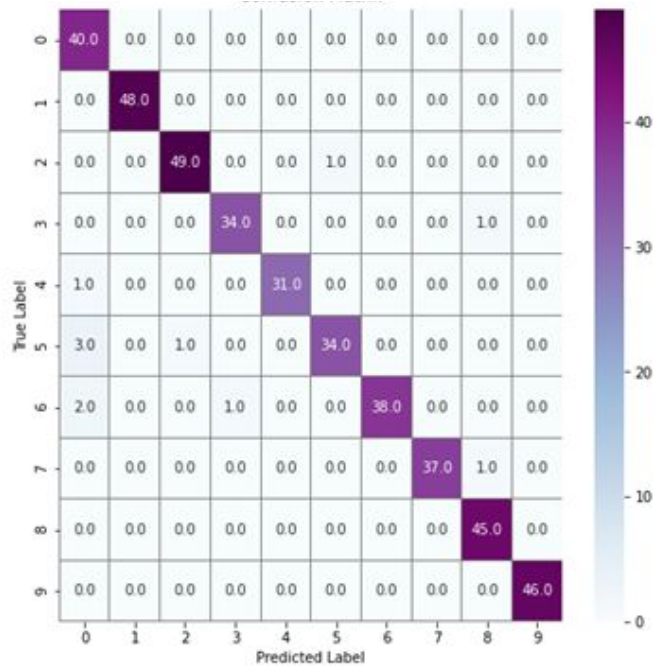
After 100 epochs, validation accuracy at 97%+ and validation loss at 11%.

Model Accuracy



Model Loss



[illegible]

Alphabet Dataset

- Dataset:
 - 34,627 images of ASL alphabet hand symbols (~1,400 per letter)
 - 28 x 28 pixel black and white images, in csv format
 - ~80/20 train vs test split
- Data transformation required:
 - Pixel values adjusted from 0-255 to between 0-1
 - One hot encoding of y class
 - Resize to 64 x 64 using Pillow function
 - Image augmentation






Image Augmentation


Original




Random Rotations




Feature Standardisation




Random Shifts



ZCA Whitening



Random Flips





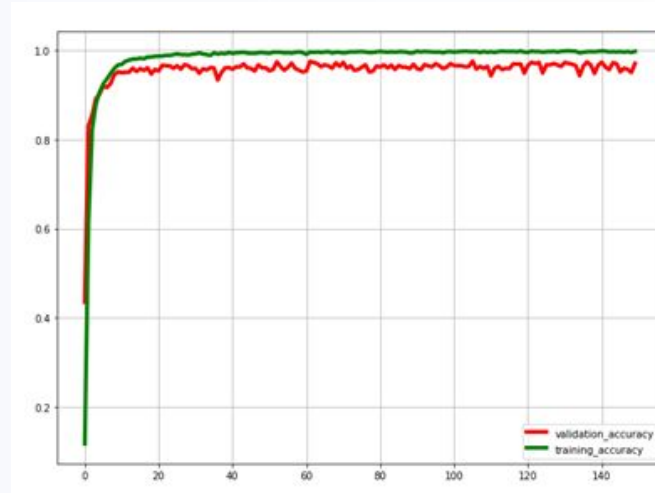
Combined Alphabet & Digits Model

- Dataset:
 - 36,684 images of ASL digit and alphabet hand symbols
 - Data amended to numpy array, using block_diag to join the 2 datasets
 - ~80/20 train vs test split
- Structure & compile used
 - Similar structure and compile used as Digits model.


Combined Alphabet & Digits Model cont.

After 150 epochs, validation accuracy at 97%+ and validation loss at 40%.

Model Accuracy




Model Loss




Combined Alphabet & Digits Model cont.

Confusion Matrix



Additional Pictures - 1

New Input



Our Expectation = 5



Model Output

its an A!



A

Additional Pictures - 2

New Input



Our Expectation = 5



Model Output

its a V!



V

Additional Pictures - 3

New Input



Our Expectation = 4



Model Output

its a Q!




Q

Additional Pictures - 4

New Input



Our Expectation = F



Model Output




Additional Pictures - 5

New Input



Our Expectation = 0



Model Output

its an 0!



0



Learnings

- Integrity of data - encountered problems with labelling / resolution of data and resulting model output
- Significant difference in model accuracy output between running on our own computer vs running on Google Colab.
- Colab run models may not be compatible to your computer. Also encountered issues with saving datasets and GPU usage.
- Reconsider flip and rotate image augmentation, these may have confused the model and created some erroneous classifications.
- Adjustment of epochs when running the model

Learnings cont.

```
] # Fit the model - on our combined dataset
# the higher the batch size, the more accurate, but also takes up more memory (esp as we are running images)
history = model3.fit(X_total_train_resized,y_total_train,epochs=150, batch_size=128, validation_data=(X_total_test_resized,y_total_test))

Epoch 39/150
228/228 [=====] - 2s 11ms/step - loss: 0.0351 - accuracy: 0.9906 - val_loss: 0.0719 - val_accuracy: 0.9779
Epoch 40/150
228/228 [=====] - 2s 10ms/step - loss: 0.0239 - accuracy: 0.9928 - val_loss: 0.0458 - val_accuracy: 0.9826
Epoch 41/150
228/228 [=====] - 2s 11ms/step - loss: 0.0237 - accuracy: 0.9932 - val_loss: 0.0557 - val_accuracy: 0.9852
Epoch 42/150
228/228 [=====] - 2s 10ms/step - loss: 0.0237 - accuracy: 0.9938 - val_loss: 0.0744 - val_accuracy: 0.9843
Epoch 43/150
228/228 [=====] - 2s 11ms/step - loss: 0.0282 - accuracy: 0.9939 - val_loss: 0.0545 - val_accuracy: 0.9835
Epoch 44/150
228/228 [=====] - 2s 10ms/step - loss: 0.0201 - accuracy: 0.9945 - val_loss: 0.0466 - val_accuracy: 0.9859
Epoch 45/150
228/228 [=====] - 2s 10ms/step - loss: 0.0157 - accuracy: 0.9950 - val_loss: 0.0714 - val_accuracy: 0.9772
Epoch 46/150
228/228 [=====] - 2s 11ms/step - loss: 0.0230 - accuracy: 0.9936 - val_loss: 0.1267 - val_accuracy: 0.9713
Epoch 47/150
228/228 [=====] - 2s 11ms/step - loss: 0.0297 - accuracy: 0.9931 - val_loss: 0.1117 - val_accuracy: 0.9760
Epoch 48/150
228/228 [=====] - 2s 10ms/step - loss: 0.0225 - accuracy: 0.9939 - val_loss: 0.1028 - val_accuracy: 0.9773
Epoch 49/150
228/228 [=====] - 2s 10ms/step - loss: 0.0192 - accuracy: 0.9938 - val_loss: 0.0856 - val_accuracy: 0.9834
Epoch 50/150
228/228 [=====] - 2s 10ms/step - loss: 0.0200 - accuracy: 0.9949 - val_loss: 0.1219 - val_accuracy: 0.9684
```




Gracias!! 🙄